

Cost estimate: workflow + Gemini vs. Claude Code

What it costs to generate a small app or website with this repo's workflow instead of asking Claude directly · 24 September 2026 · Companion spreadsheet: cost-estimate.xlsx (all figures below are its formulas)

Summary

- **The workflow is about 10× cheaper per job.** A static website costs \$0.026 on paid Gemini 3.6 Flash (and \$0 on the free tier until 31 Dec 2026), versus roughly \$0.25 (Sonnet 5) to \$0.44 (Opus 5.5) for a typical Claude Code session billed through the API.
- **Both are cheap at low volume.** The gap only matters at scale: 100 sites a month is \$2.58 with the workflow versus about \$24.90 with Claude Code on Sonnet 5.
- **Cost per call is not cost per working result.** The workflow writes a site nobody has opened or tested; Claude Code checks and fixes its own work. Any manual rework quickly outweighs the API difference.
- **On a Claude Pro or Max plan, Claude Code has no per-job cost** within your usage limits (\$20 or \$100–200 a month, flat).

1. Prices used (USD per 1M tokens)

Model	Input	Output	Cache read	Cache write	Notes
Gemini 3.6 Flash (free tier)	\$0	\$0	–	–	Free until 31 Dec 2026; rate-limited
Gemini 3.6 Flash (paid, 2026)	\$0.75	\$3.75	–	–	Paid tier, 2026
Gemini 3.6 Flash (paid, from 2027)	\$1.50	\$7.50	–	–	Paid tier from 1 Jan 2027
Claude Sonnet 5	\$2.00	\$10.00	\$0.20	\$2.50	Cache read/write assumed 0.1× / 1.25×
Claude Opus 5.5	\$4.00	\$20.00	\$0.20	\$5.00	Cache write assumed 1.25×

Subscriptions (include Claude Code, with usage limits): Claude Pro \$20/month · Max 5× \$100/month · Max 20× \$200/month.

2. Tokens per job

Approach	Job	Calls / turns	Input tokens	of which cached	Output tokens	Basis
Workflow	Static website	1	295	–	6,814	Measured
Workflow	Tested backend - passes first try	1	333	–	2,407	Measured
Workflow	Tested backend - worst case	4	7,233	–	9,628	Estimate
Claude Code (typical)	Static website	12	300,000	270,000	12,000	Estimate
Claude Code (heavy)	Static website	25	1,000,000	900,000	30,000	Estimate
Claude Code (typical)	Tested backend	15	450,000	405,000	15,000	Estimate
Claude Code (heavy)	Tested backend	30	1,350,000	1,215,000	35,000	Estimate

Measured rows (shaded) are real Gemini 3.6 Flash calls from this repo's workflows. Worst case = first call plus 3 failed retries, each resending the code and test output (~2,300 input tokens). Claude Code rows are ballpark estimates: turns × context size, 90% served from the prompt cache.

3. Cost per job (USD)

Approach	Job	Gemini free	Gemini paid 2026	Gemini paid 2027	Claude Sonnet 5	Claude Opus 5.5
Workflow	Static website	\$0	\$0.026	\$0.052	\$0.069	\$0.137
Workflow	Tested backend - passes first try	\$0	\$0.009	\$0.019	\$0.025	\$0.049
Workflow	Tested backend - worst case	\$0	\$0.042	\$0.083	\$0.111	\$0.221
Claude Code (typical)	Static website	n/a	n/a	n/a	\$0.249	\$0.444
Claude Code (heavy)	Static website	n/a	n/a	n/a	\$0.730	\$1.280
Claude Code (typical)	Tested backend	n/a	n/a	n/a	\$0.344	\$0.606
Claude Code (heavy)	Tested backend	n/a	n/a	n/a	\$0.930	\$1.618

"n/a": Claude Code only runs Claude models. The Claude columns on workflow rows show what the same n8n workflow would cost on the Claude API (tokenizers differ, so approximate).

4. Monthly cost — static website

Option	Per job	10 jobs / month	100 jobs / month	1,000 jobs / month
Workflow + Gemini 3.6 Flash (free tier)	\$0	\$0	\$0	\$0
Workflow + Gemini 3.6 Flash (paid, 2026)	\$0.026	\$0.26	\$2.58	\$25.77
Workflow + Gemini 3.6 Flash (paid, from 2027)	\$0.052	\$0.52	\$5.15	\$51.55
Workflow + Claude Sonnet 5 API	\$0.069	\$0.69	\$6.87	\$68.73
Claude Code, API billing, Sonnet 5 (typical)	\$0.249	\$2.49	\$24.90	\$249.00
Claude Code, API billing, Opus 5.5 (typical)	\$0.444	\$4.44	\$44.40	\$444.00
Claude Pro subscription (flat, usage limits apply)	flat	\$20.00	\$20.00	\$20.00
Claude Max 5x subscription (flat)	flat	\$100.00	\$100.00	\$100.00

Subscriptions are flat but capped: 1,000 Claude Code sessions a month would exceed a Pro plan's limits. The free Gemini tier is also rate-limited, so high volumes may need the paid tier.

5. Which to use

- **Use the workflow + Gemini** for many cheap, unattended, "good-enough" drafts — triggered by a form, a schedule or another system — where a person reviews the output afterwards.
- **Use Claude Code** when you want one site or app that is known to work: it reads your repo, runs the code, and fixes problems in the same session. It costs more per job, but typically well under a few dollars.

Caveats

- Claude Code token counts depend heavily on the task and the number of fix rounds; edit the assumptions in the spreadsheet's *Token Usage* sheet to match your own sessions.
- Gemini 3.6 Flash prices double on 1 January 2027, and its free tier ends on 31 December 2026. Review Google's free-tier data terms before sending anything private.
- Sonnet 5 cache-read and Claude cache-write prices use the standard 0.1× / 1.25× multipliers (assumed).
- Running n8n and Docker on your own machine adds no direct cost.

Sources

Gemini API pricing — <https://ai.google.dev/gemini-api/docs/pricing> · Claude plans & pricing — <https://claude.com/pricing> · Claude API model prices — Anthropic Claude API reference (cached 24 June 2026) · Claude Code pricing overview — <https://duet.so/blog/claude-code-pricing> · Token measurements — real Gemini 3.6 Flash calls made from this repo's workflows, 24 September 2026